



Chapter 1:Introduction to Machine Learning

Ph.D. Nguyen Thi Khanh Tien
tienntk@ut.edu.vn

DS Roadmap (Recap)

Data Scientist

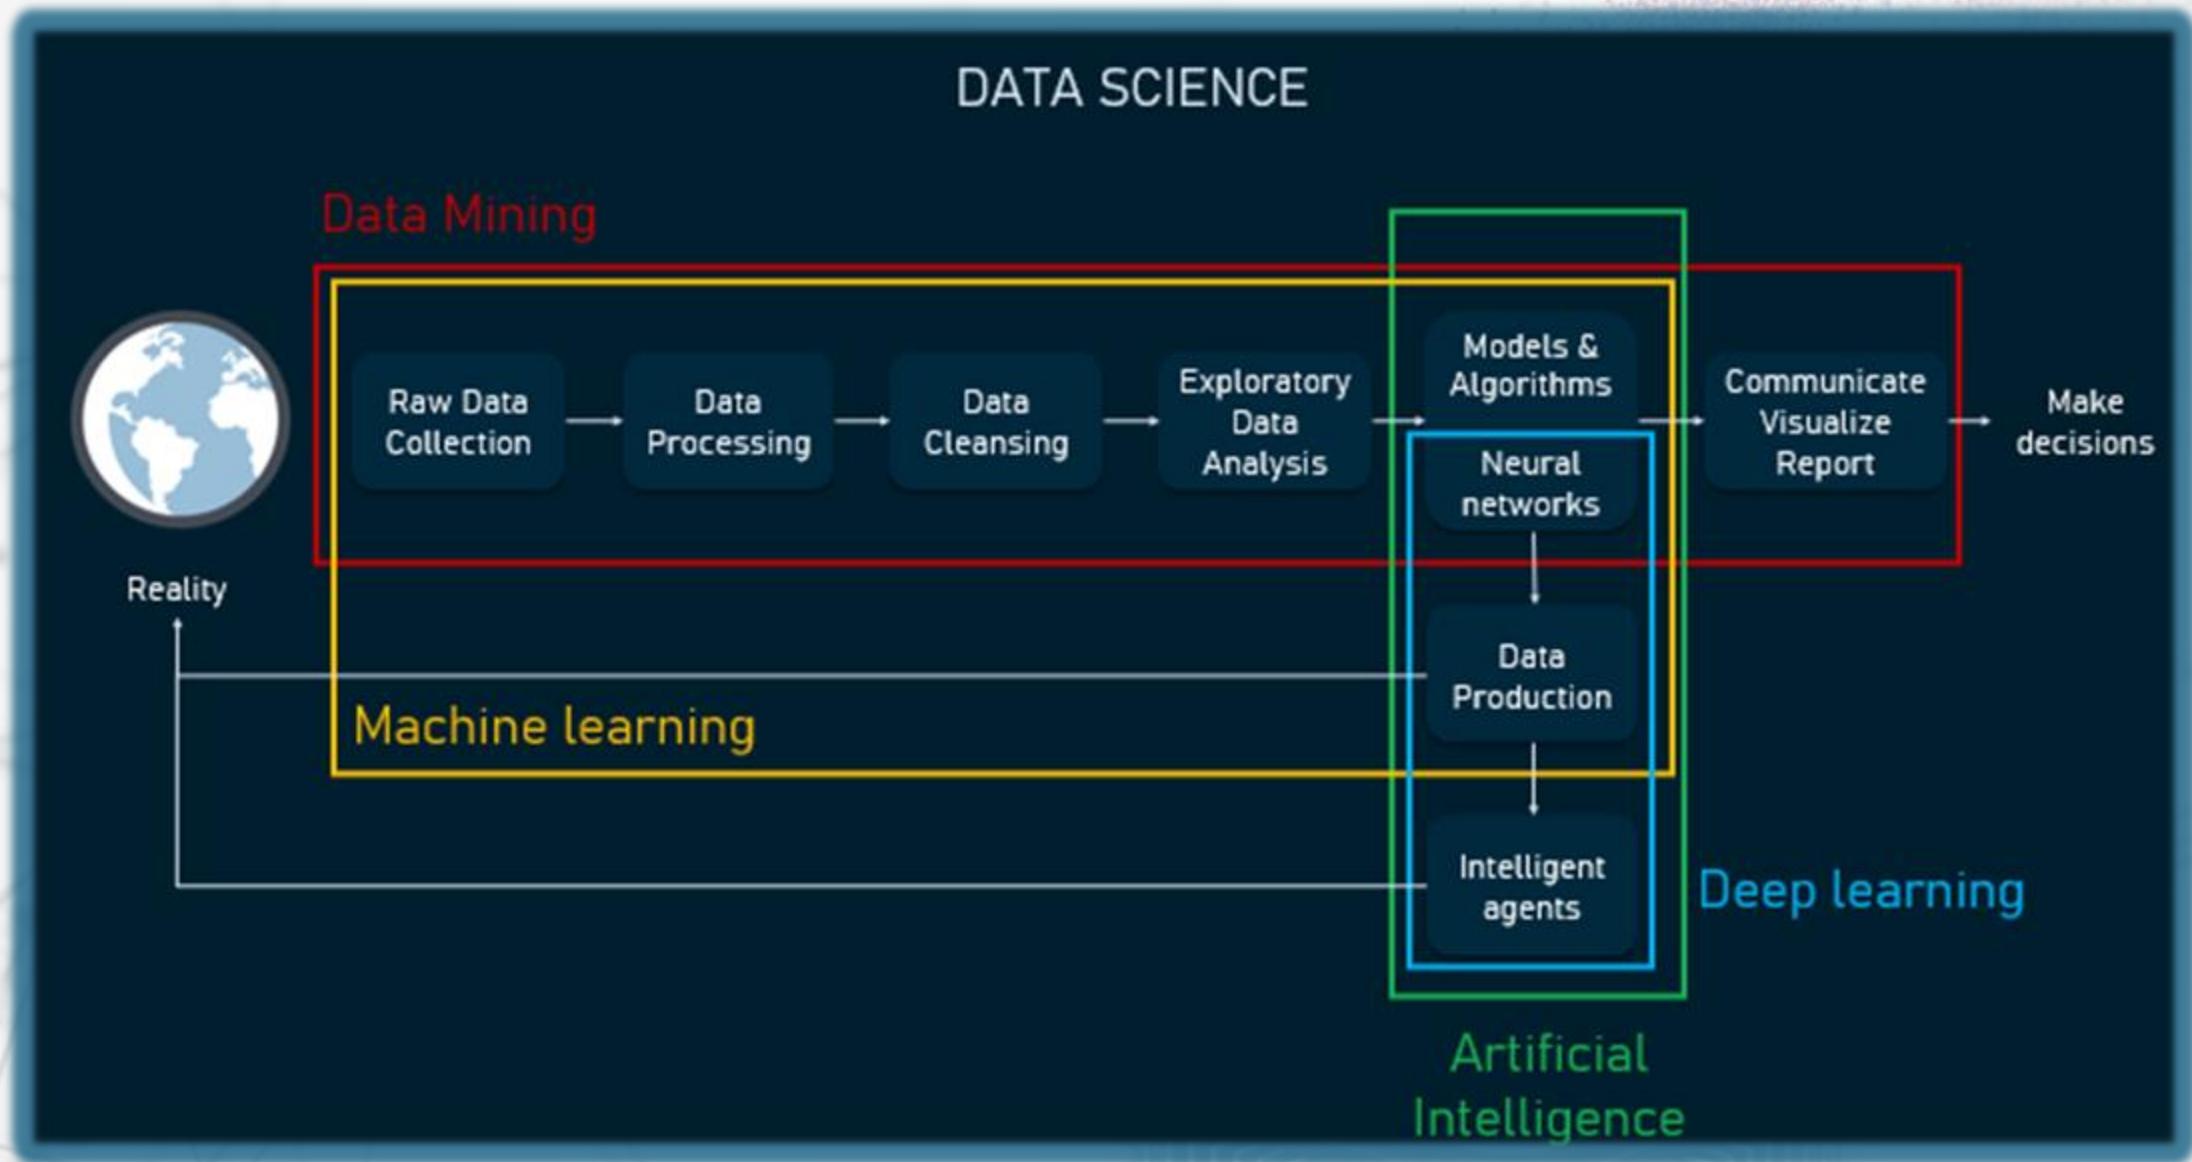
Roadmap



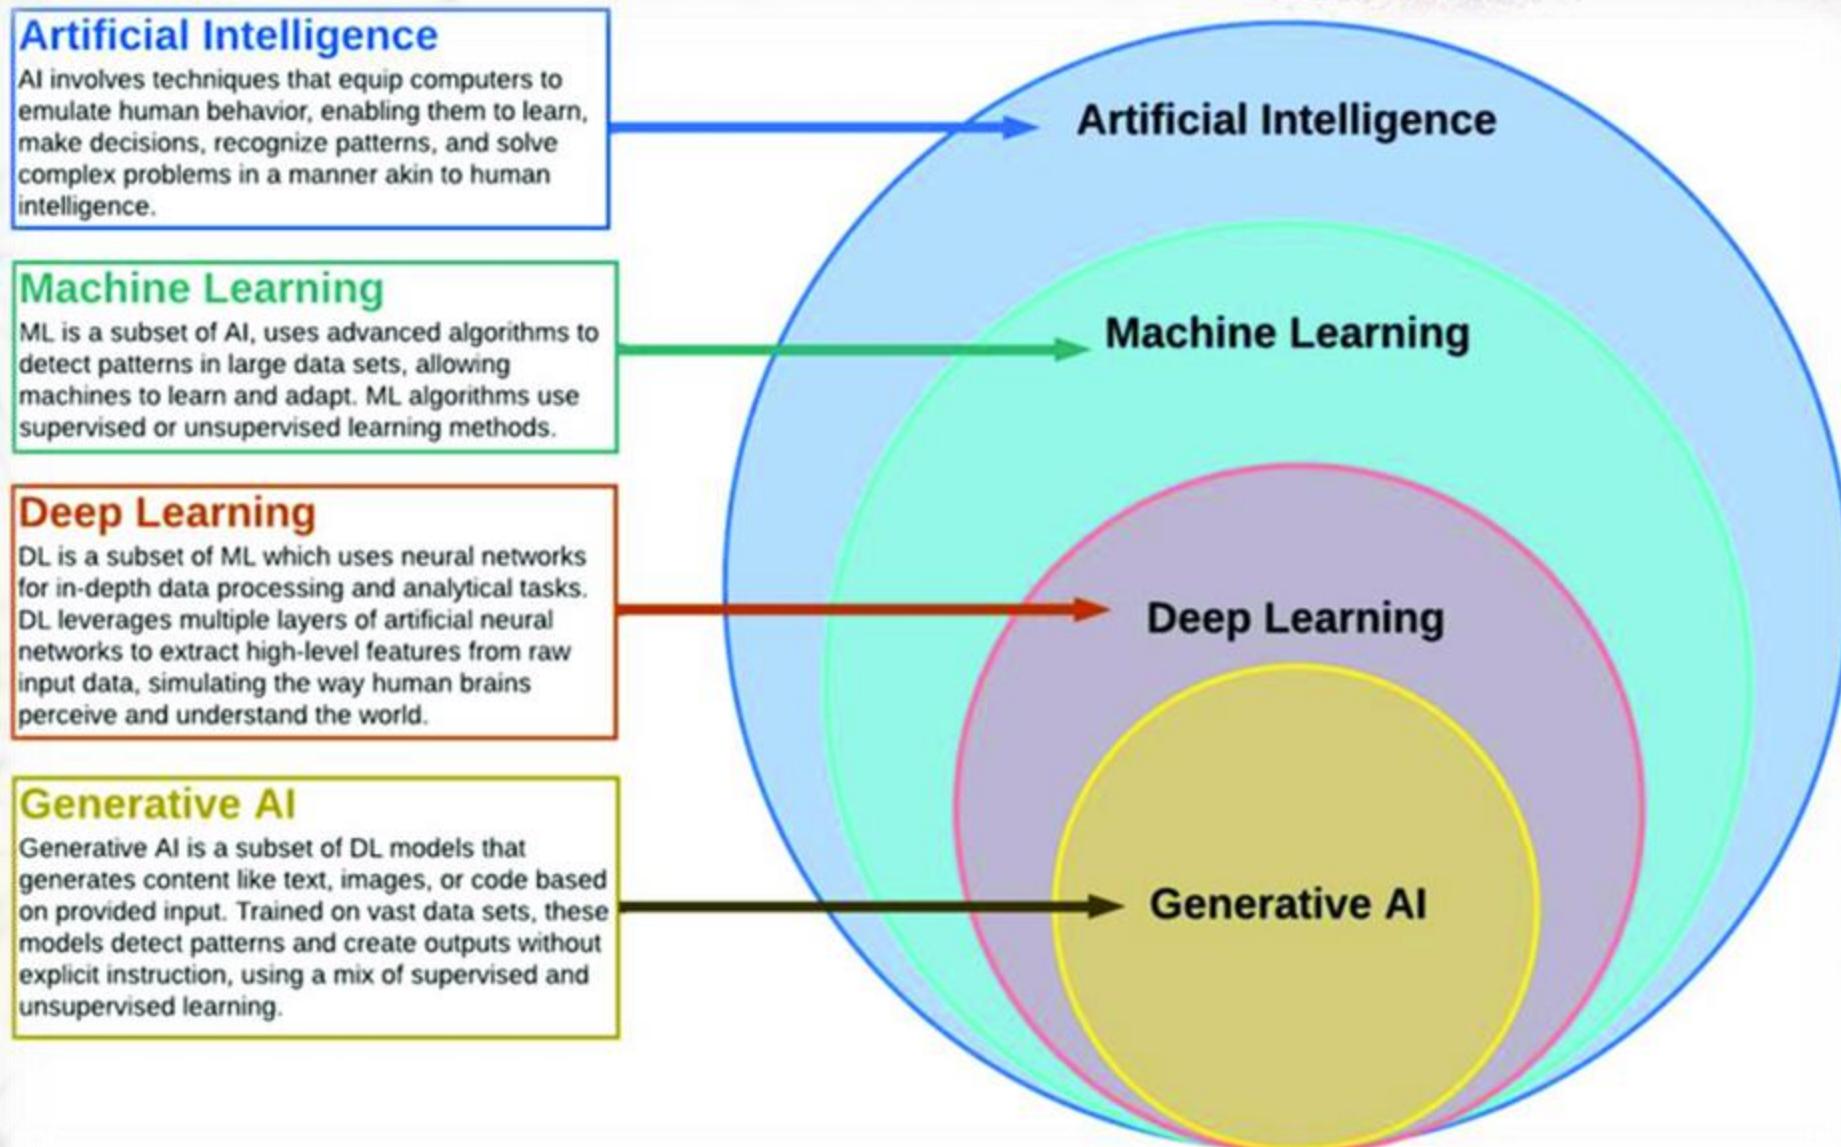
DS Roadmap
(Recap)



Data Science Pipeline



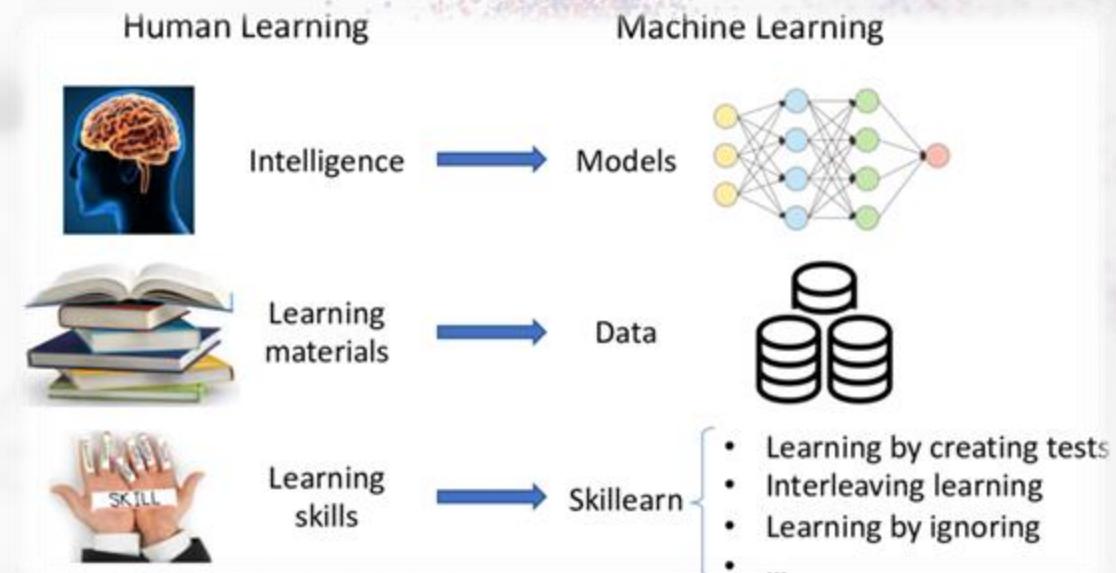
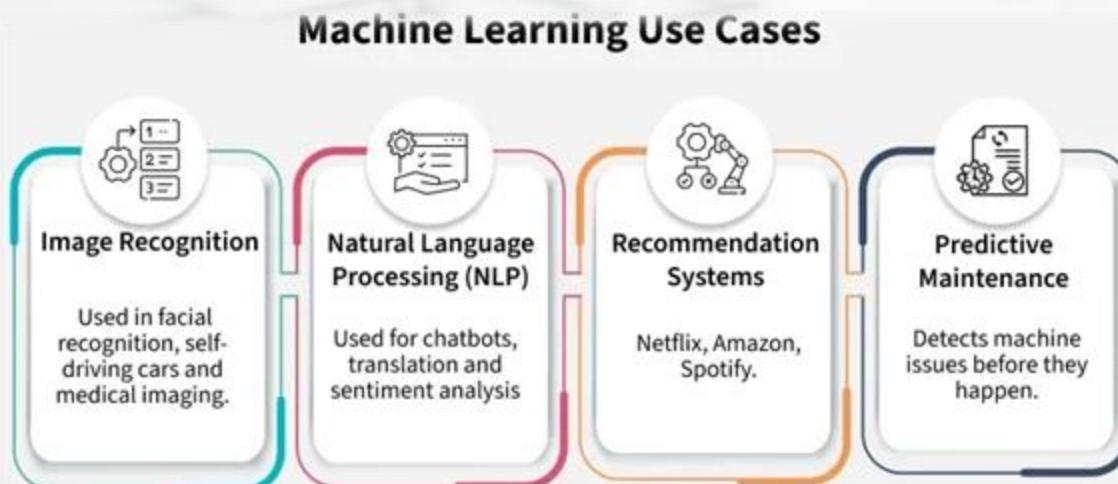
AI, ML, DL, GenAI



Understanding Machine Learning

Machine learning is a subset of artificial intelligence that enables computers to learn from data and improve their performance on a specific task without being explicitly programmed.

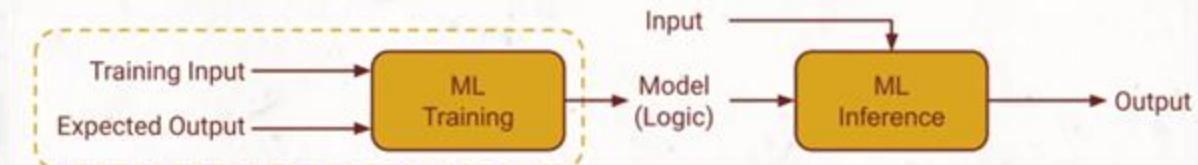
Python has become the de facto language for machine learning due to its simplicity, readability, and extensive libraries.



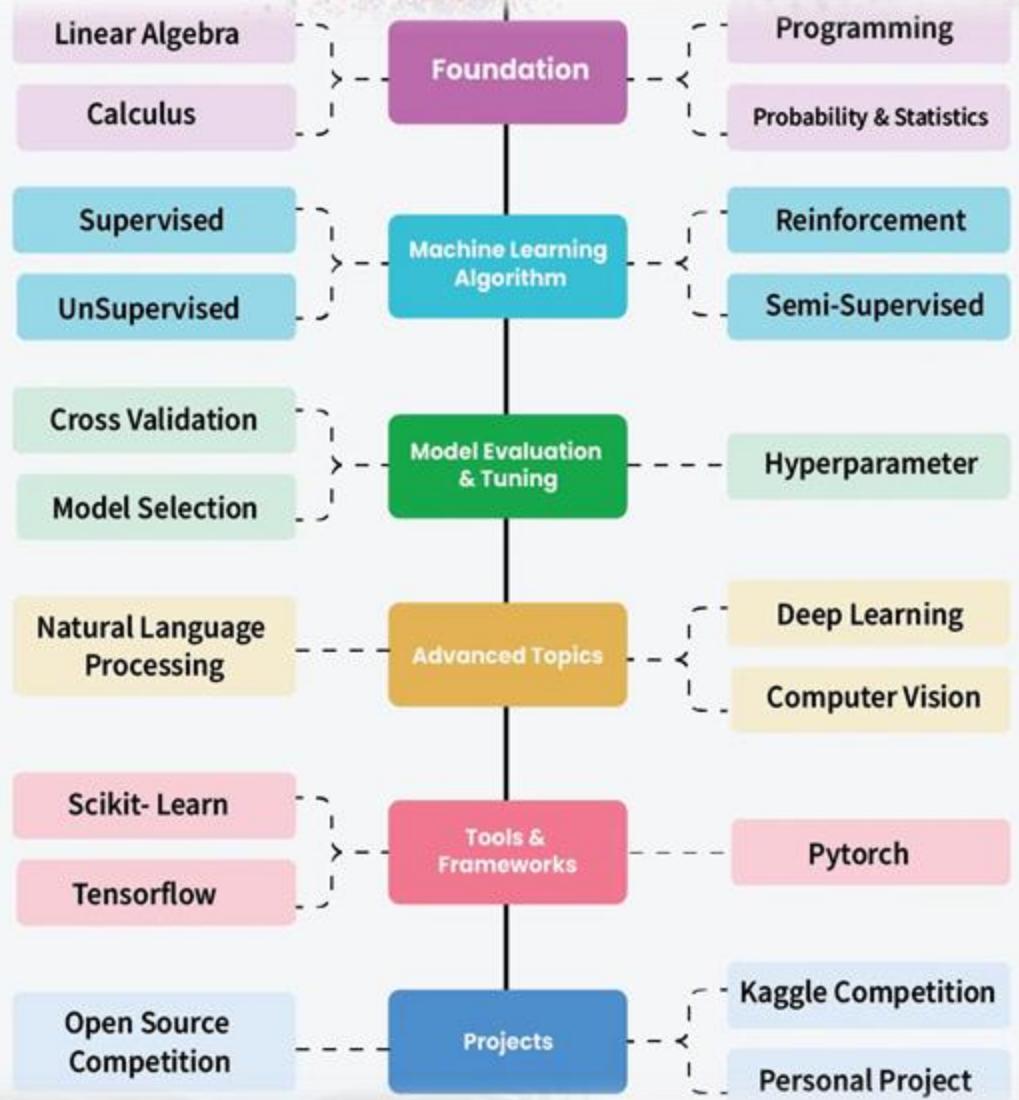
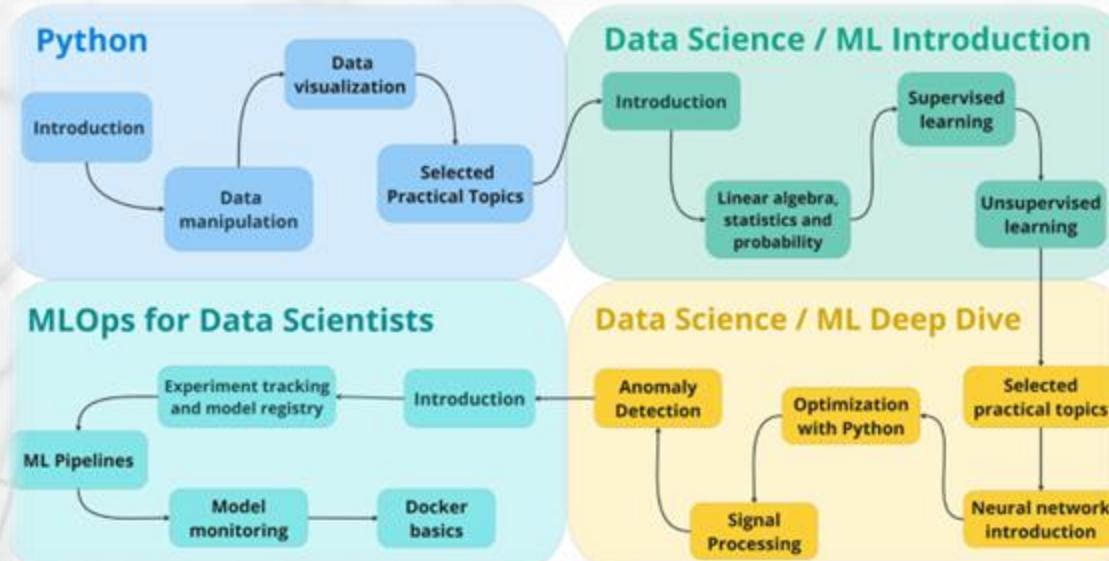
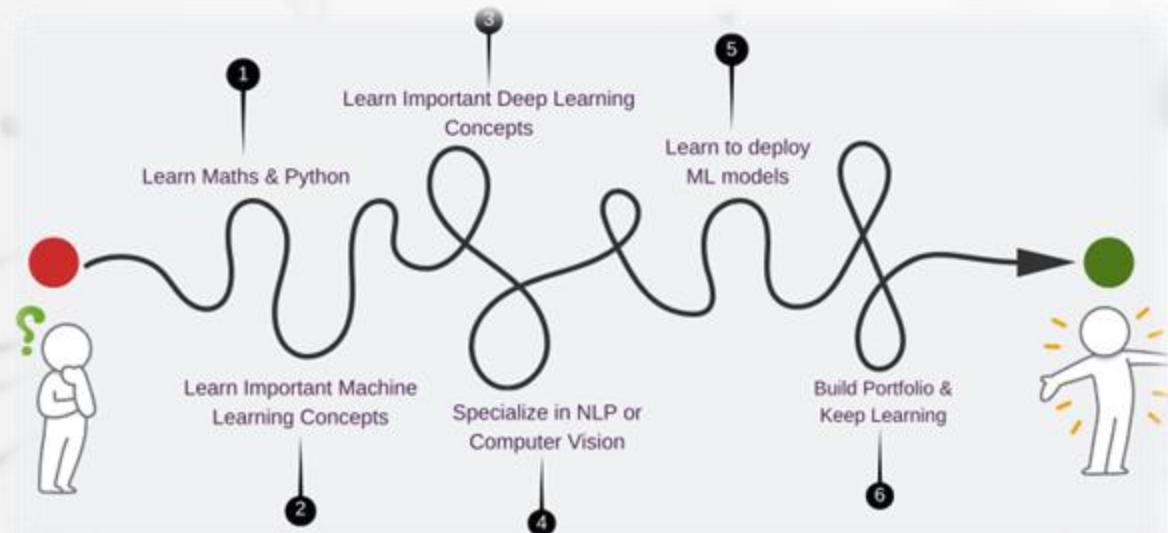
Traditional Programs: Define algo/logic to compute output



Machine Learning: Learn model/logic from data



Roadmap to learn ML

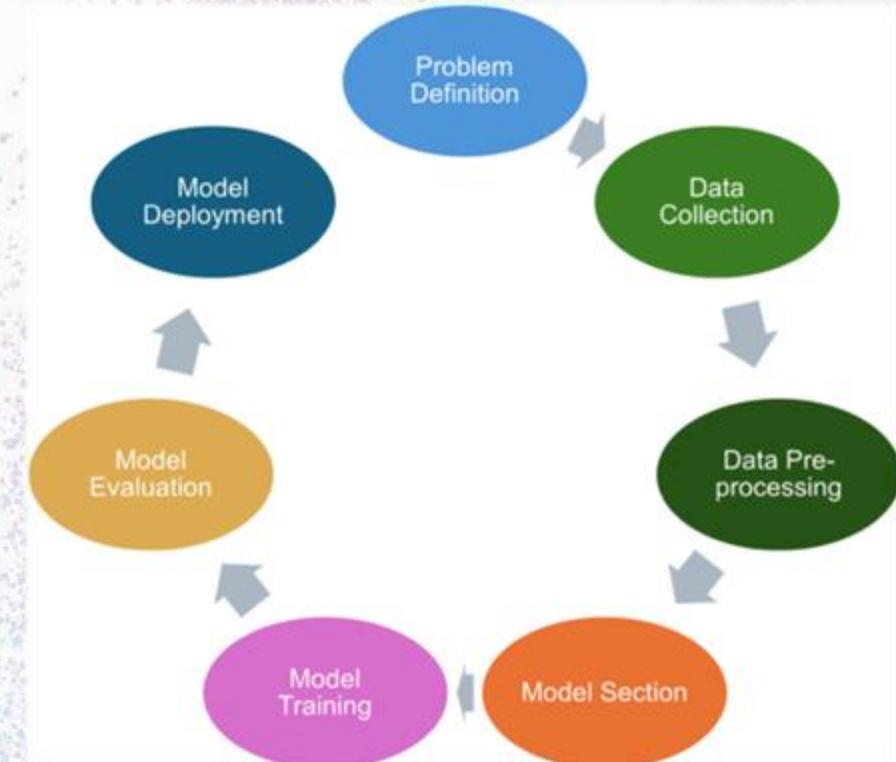


Roadmap to learn ML

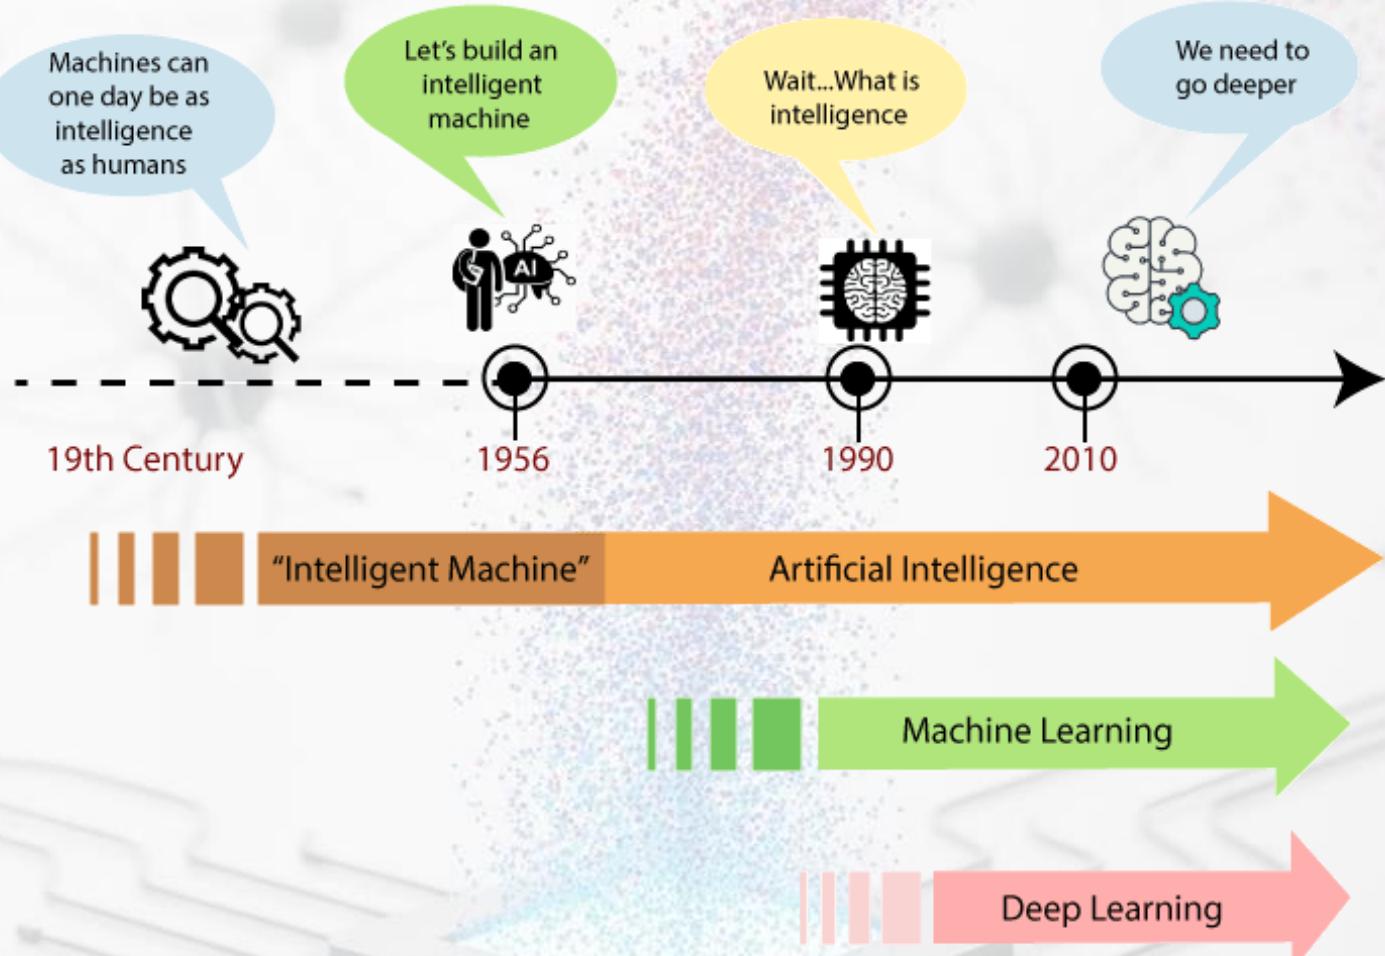
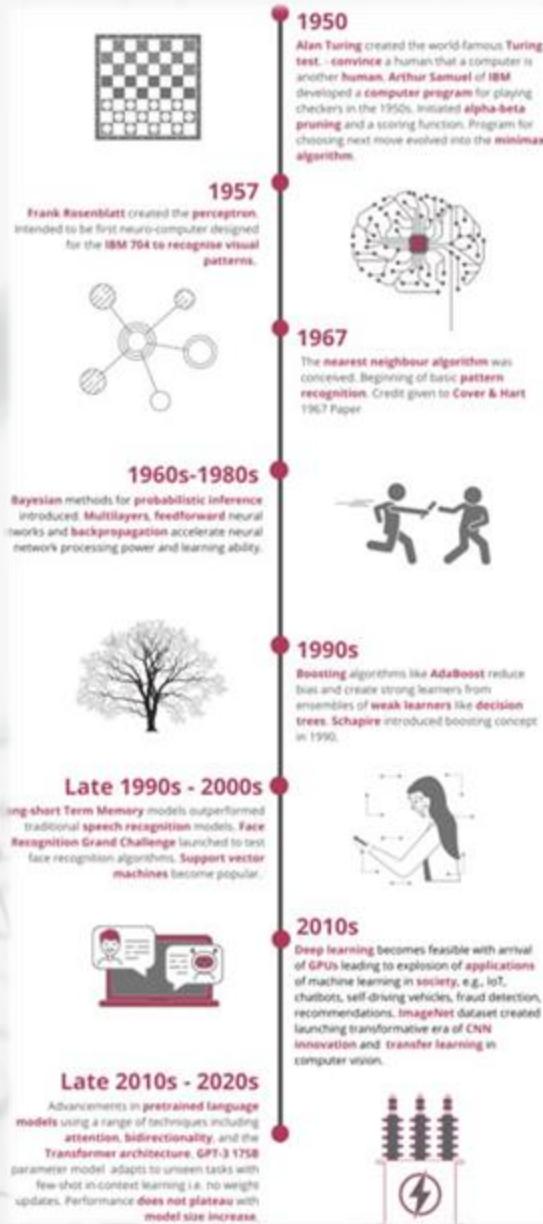


Key Concepts in Machine Learning

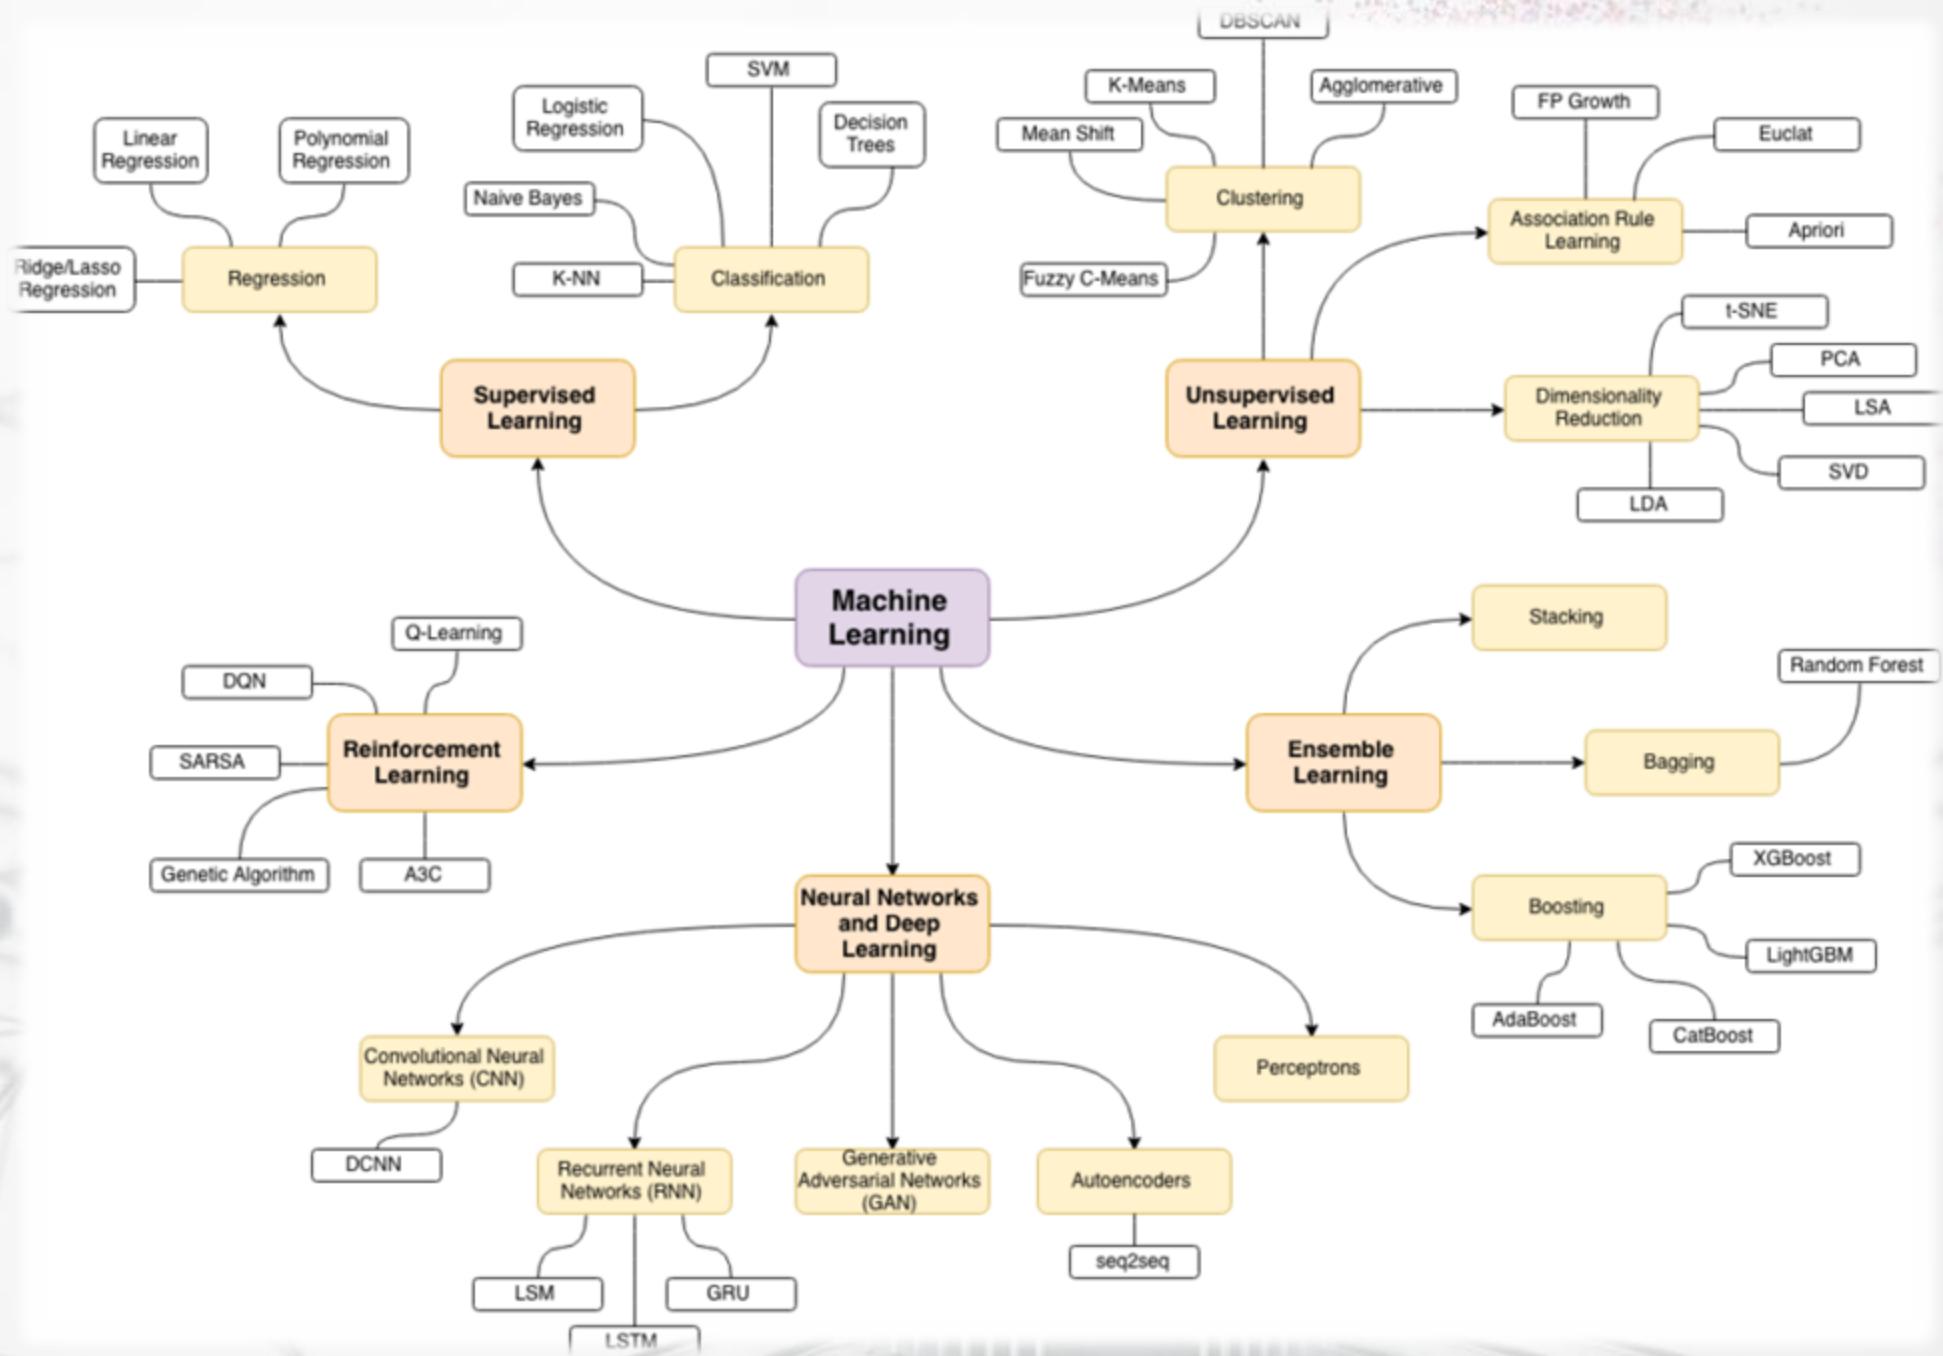
- **Data:** The raw material that machines learn from. It can be structured (e.g., CSV files) or unstructured (e.g., images, text).
- **Algorithms:** The methods used to process data and extract patterns. Examples include linear regression, decision trees, random forests, and neural networks.
- **Model:** A representation of the learned patterns. It can be used to make predictions on new, unseen data.
- **Training:** The process of feeding data to an algorithm to teach it to recognize patterns.
- **Testing:** Evaluating the model's performance on a separate dataset to assess its accuracy.



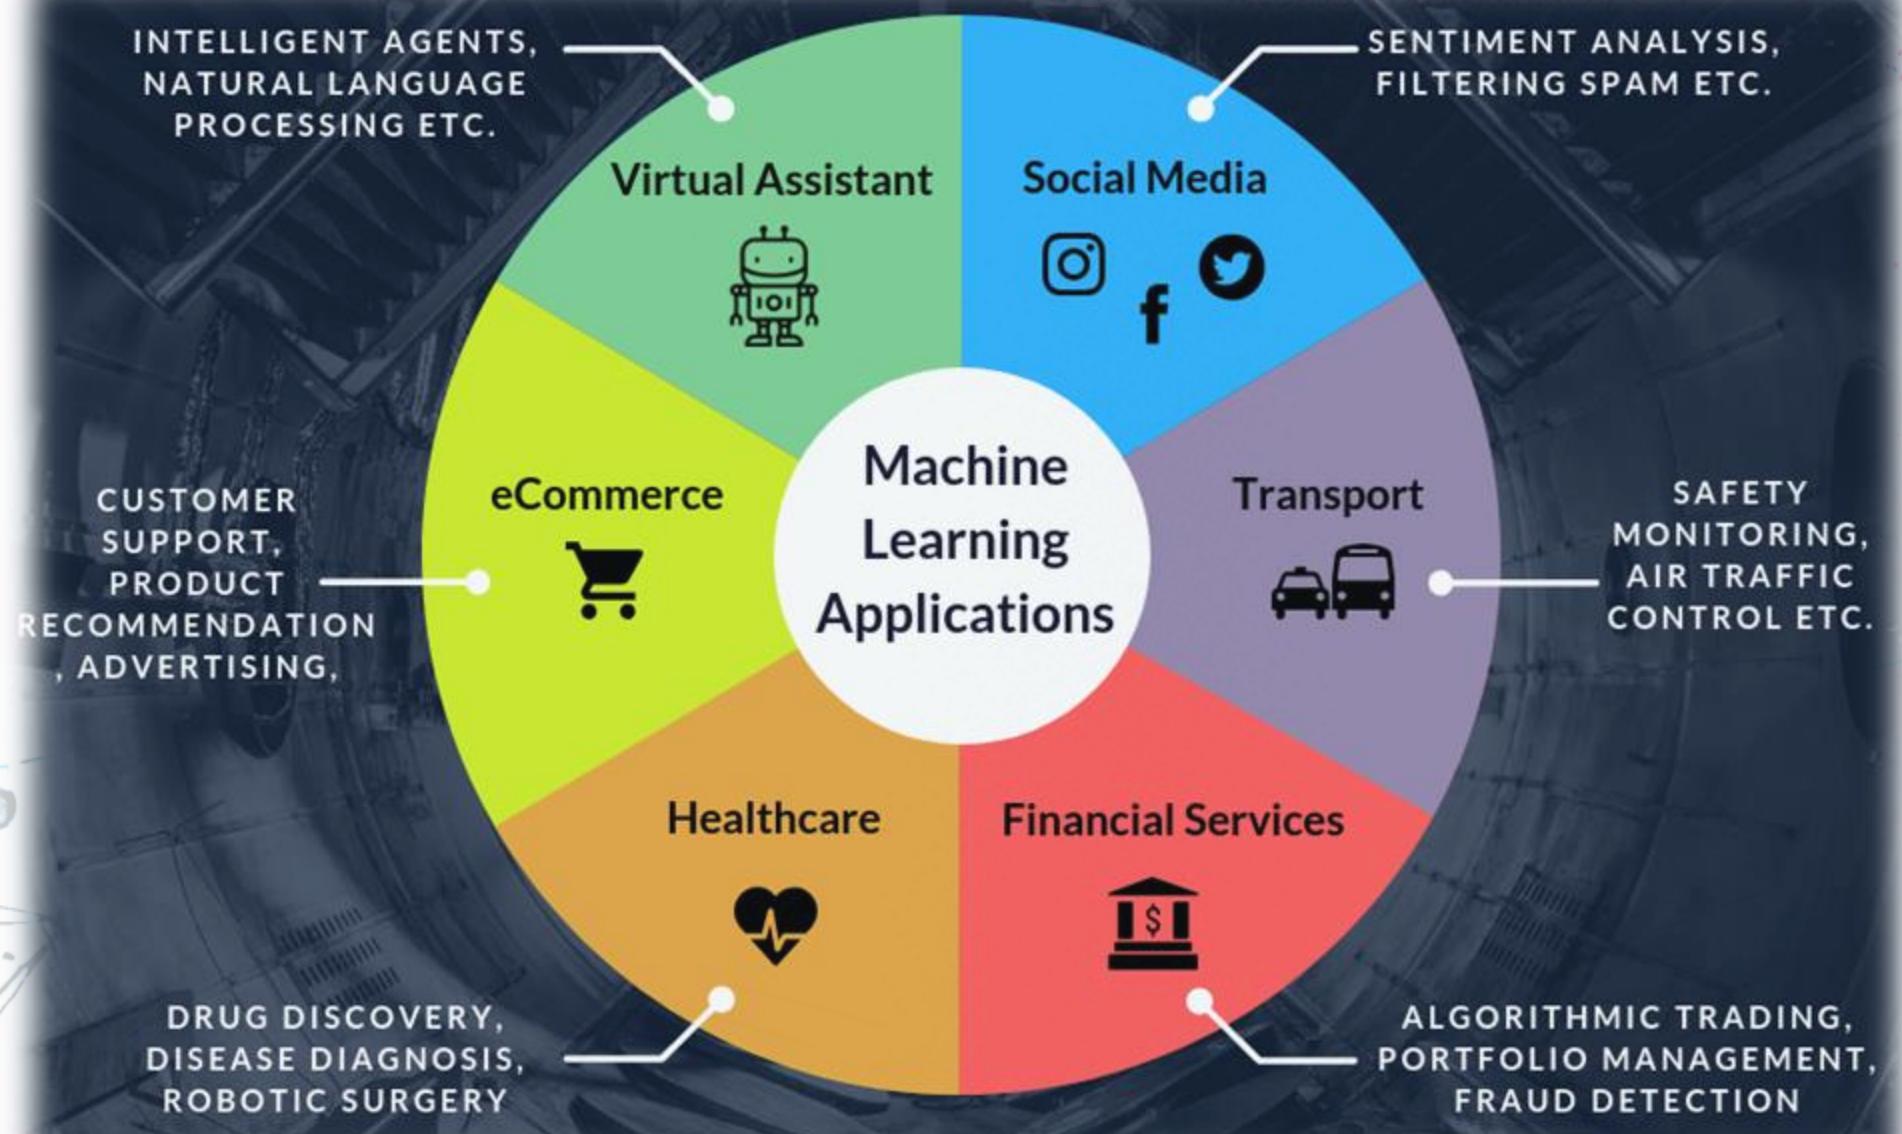
A Brief History of ML



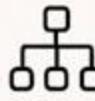
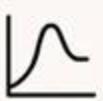
The map of ML



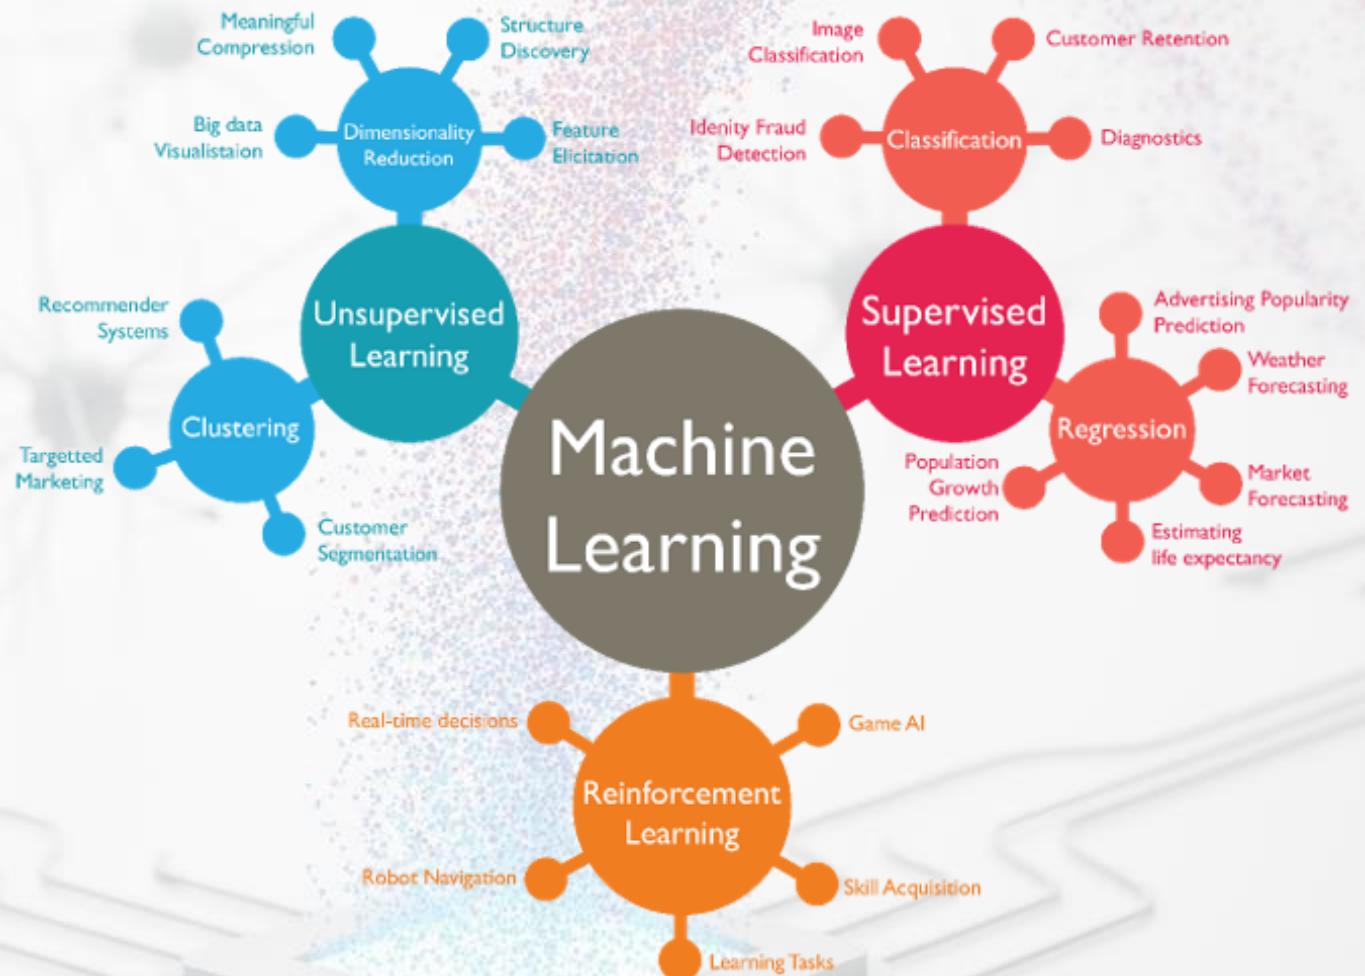
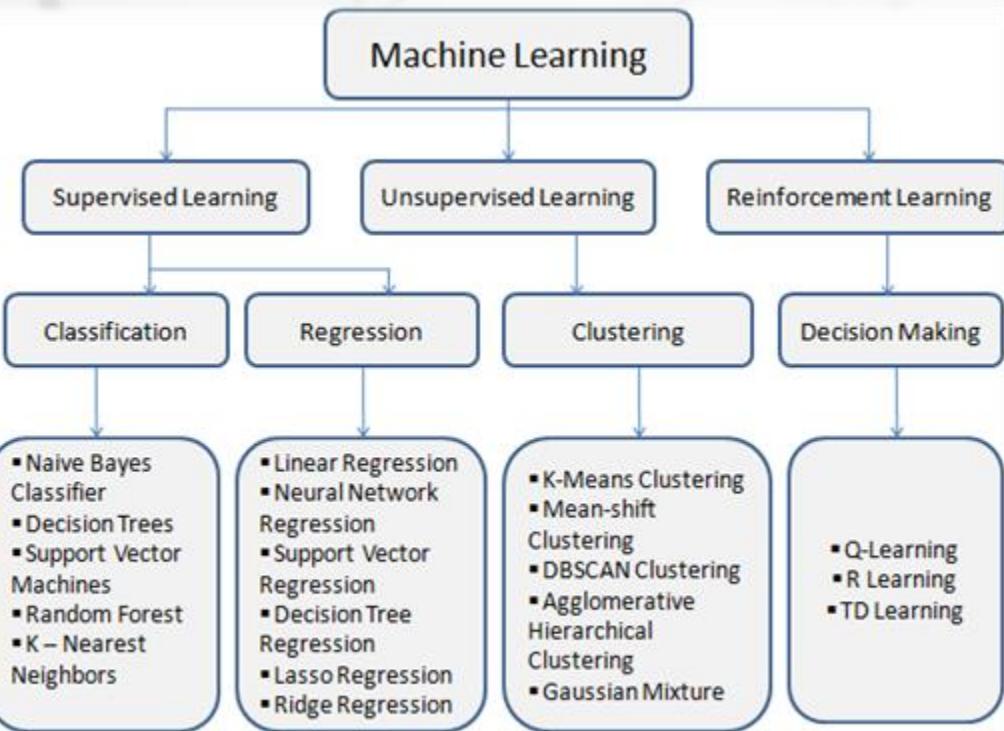
Application of ML



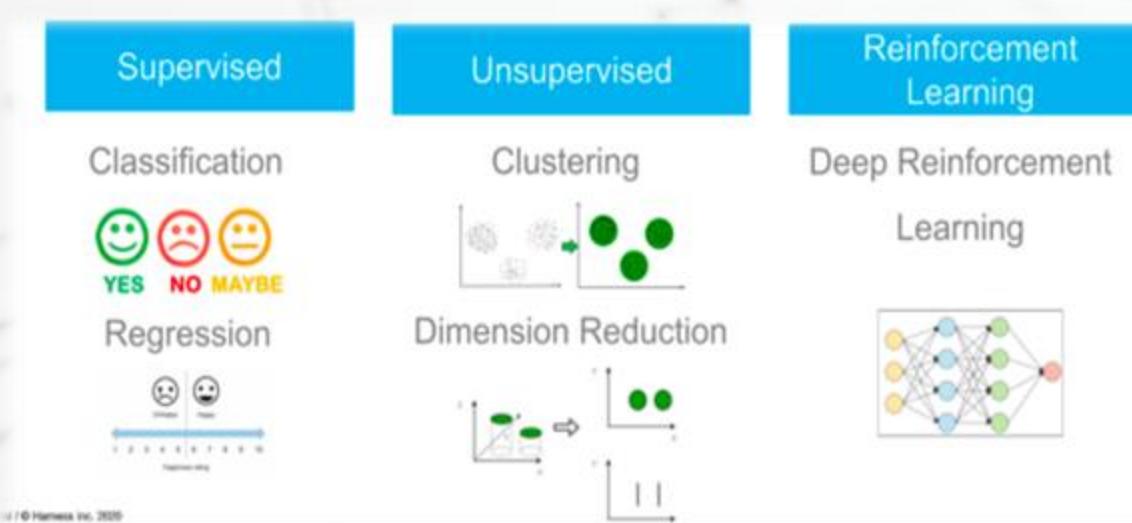
Advantage & Disadvantage of ML

 Advantages	 Disadvantages
 Automation ML can automate repetitive tasks and decision-making processes.	 Data Dependency Needs large, quality datasets to perform well.
 Self-Learning Improves performance over time as more data becomes available.	 Training Time Some models (like deep learning) can be computationally expensive.
 Handles Complex Data Excels at uncovering patterns in large, high-dimensional datasets	 Interpretability Models like neural networks are often black-box and hard to explain.
 Wide Application Used in healthcare, finance, recommendation systems, NLP, etc.	 Bias & Fairness ML can inherit bias from data, leading to unfair or unethical results
 Real-Time Processing Enables predictions and decisions in real	 Overfitting If not properly validated, models can memorize data instead of generalizing

Type of ML



Type of ML

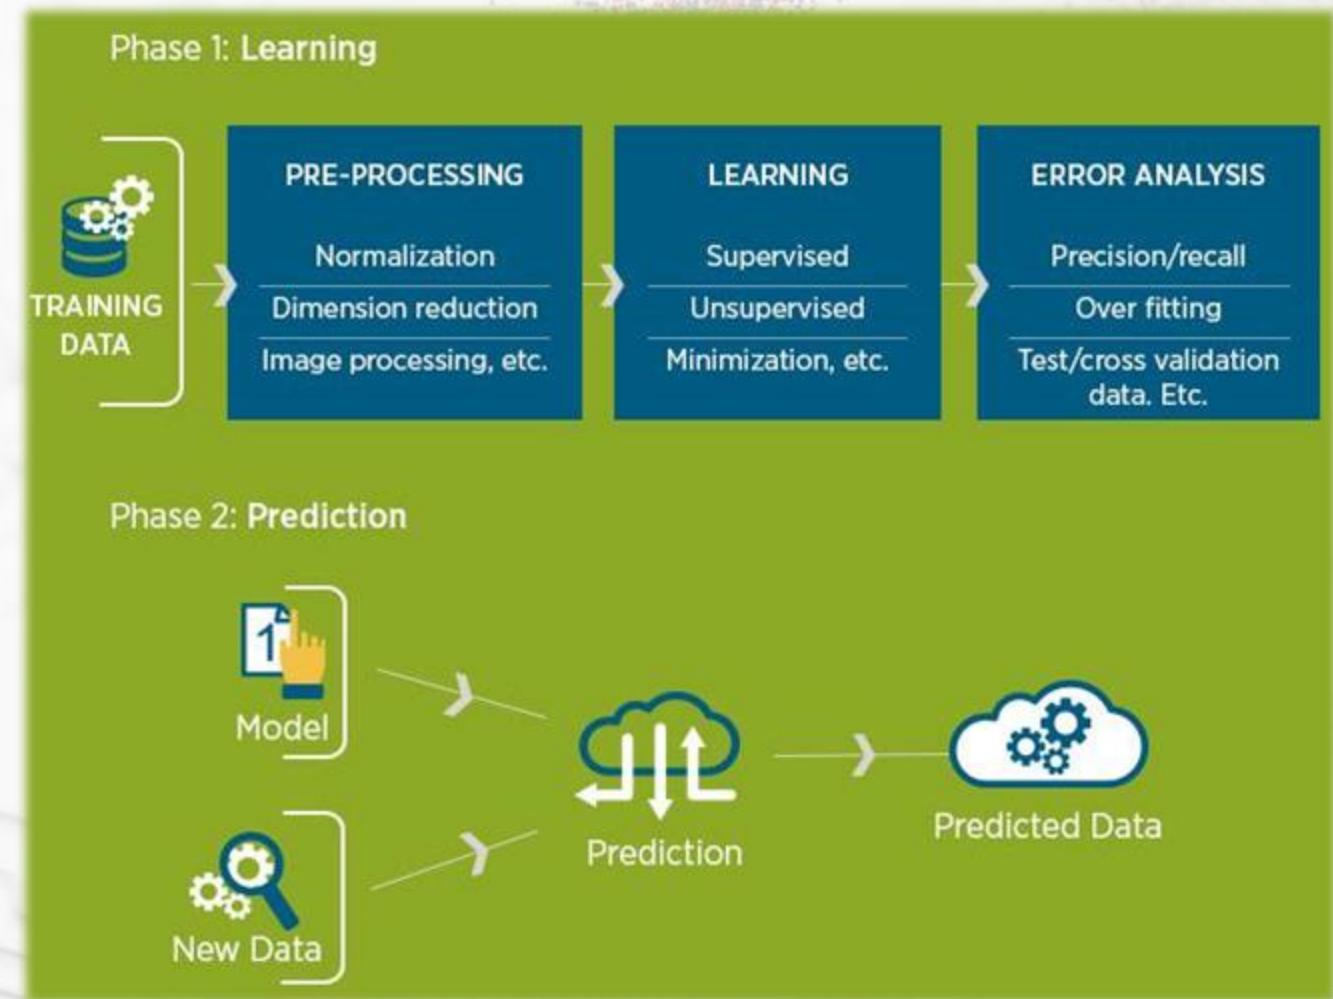
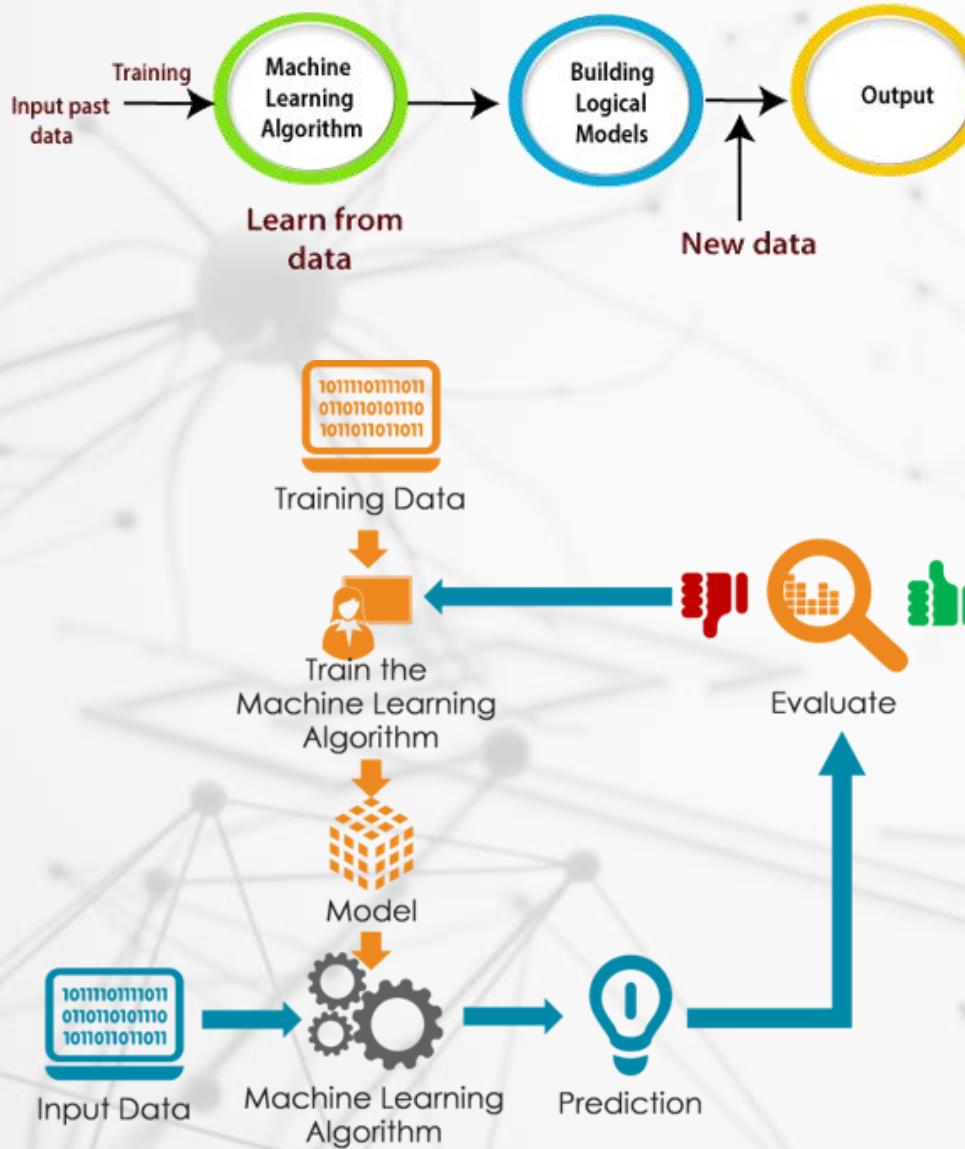


Which to Use When?

- **Supervised Learning:** When you have labeled data and want to make predictions or classifications.
- **Unsupervised Learning:** When you have unlabeled data and want to discover hidden patterns or group similar data points.
- **Reinforcement Learning:** When you want to train an agent to make decisions in an environment to maximize rewards.

Feature	Supervised Learning	Unsupervised Learning	Reinforcement Learning
Data	Labeled data	Unlabeled data	No predefined data, learns from interactions with the environment
Goal	Predict outputs based on inputs	Discover patterns and structures	Learn a policy to maximize rewards
Supervision	Requires human supervision	No human supervision	No explicit supervision, learns through trial and error

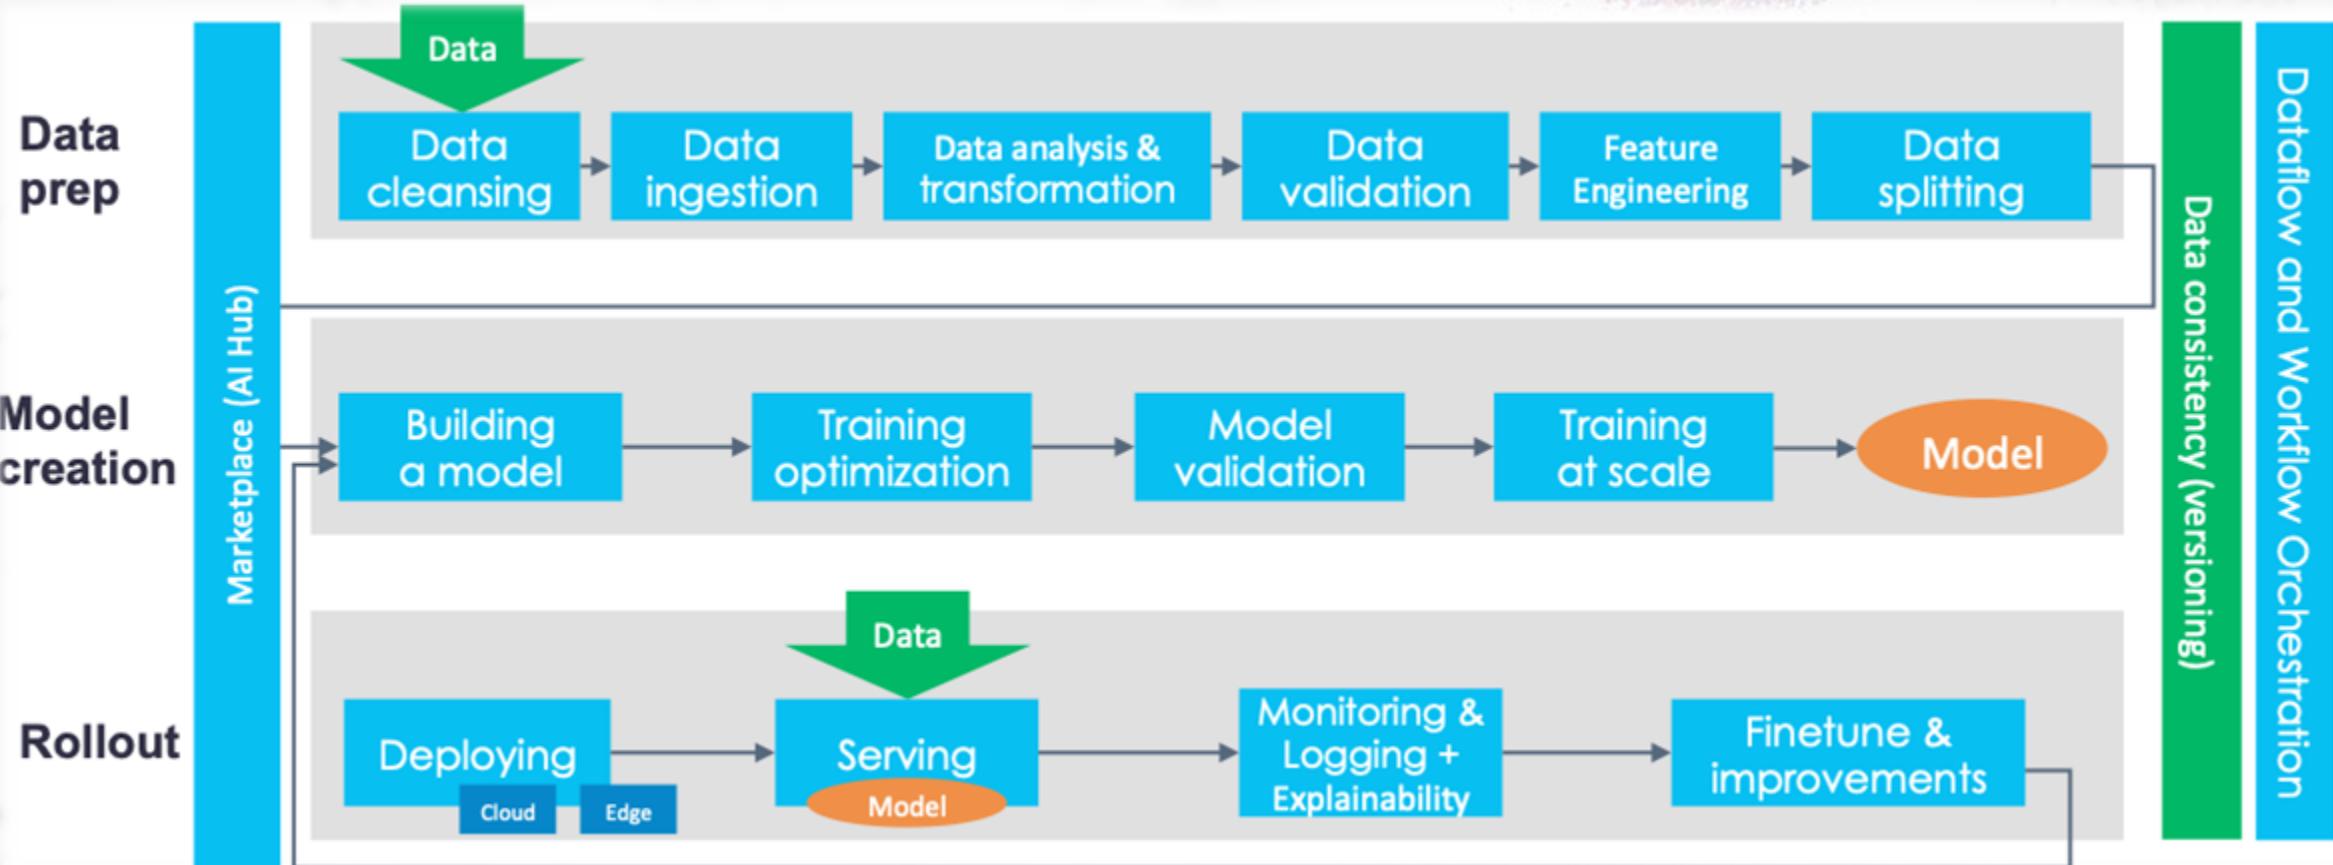
ML process



Workflow of a Machine Learning Project



Workflow of a Machine Learning Project



Basic Workflow of a Machine Learning Project

- Data Collection:** Collect relevant and high-quality data aligned with the problem you're solving.
- Data Preparation:** Clean, preprocess, and transform the raw data into a usable format (handle missing values, encoding, normalization, etc.).
- Feature Engineering:** Create or select key features that improve the model's ability to learn patterns.
- Model Selection:** Choose the most suitable algorithm based on your task (classification, regression, etc.) and data characteristics.
- Model Training:** Fit the selected model to your training dataset to learn underlying patterns.
- Model Evaluation:** Test model performance using appropriate metrics (e.g., accuracy, F1-score, MSE) on validation/test data.
- Model Deployment:** Deploy the trained model into a real-world application or system for live predictions.

Example: Predicting House Prices

- Data Collection:** Obtain a dataset with house prices and related attributes (e.g., area, location, bedrooms).
- Data Preparation:** Handle missing data, encode categorical variables, normalize numerical values.
- Feature Engineering:** Add derived features (e.g., price per square meter, age of the house).
- Model Selection:** Use regression models such as Linear Regression or Random Forest.
- Model Training:** Train on historical house data.
- Model Evaluation:** Use MSE or R² to evaluate prediction accuracy.
- Model Deployment:** Integrate into a web app or real estate tool for pricing estimation.

Important AI / ML Python Packages



→ <https://docs.vultr.com/deep-dive-into-important-ai-ml-python-packages#usability-tools>

Key Libraries of Python used for ML			
Category	Library	Primary Purpose	Key Features & Use Cases
Foundational / Data Ops	NumPy	Numerical computing with multi-dimensional arrays	Efficient array operations; mathematical functions; foundation for other ML libraries; data representation (vectors, matrices, tensors).
	Pandas	Data manipulation and analysis (tabular data)	DataFrames for structured data; data loading, cleaning, transformation; handling missing values; merging/joining datasets; essential for data preprocessing.
	Matplotlib	Static, animated, and interactive data visualization	Highly customizable plots (line, bar, scatter, etc.); essential for EDA and presenting results.
	Seaborn	High-level statistical data visualization (built on Matplotlib)	Simplifies complex statistical plots (heatmaps, violin plots, pair plots); enhances Matplotlib's aesthetics for statistical graphics.
	SciPy	Scientific and technical computing (advanced math functions)	Optimization, linear algebra, integration, signal processing; often used by other ML libraries under the hood or for specific pre-processing.
Core Machine Learning	Scikit-learn	Traditional Machine Learning algorithms and tools	Comprehensive range of supervised (classification, regression) and unsupervised (clustering, dimensionality reduction) algorithms; consistent API; includes tools for data preprocessing, model selection, and evaluation (e.g., cross-validation, metrics).
Deep Learning Frameworks	TensorFlow	End-to-end open-source platform for deep learning	High scalability; powerful for building and deploying large-scale neural networks (CNNs, RNNs, Transformers); strong for production environments; developed by Google.
	PyTorch	Open-source deep learning framework with dynamic computation graphs	Flexible and "Pythonic" API; favored in research for rapid prototyping and experimentation; easier debugging; strong community support; developed by Facebook AI Research.
	Keras	High-level Neural Networks API (runs on TensorFlow, PyTorch, JAX)	Simplifies deep learning model building and training; designed for fast experimentation; very beginner-friendly; provides a simpler interface to complex deep learning frameworks.
Specialized / Boosting	XGBoost	Optimized Gradient Boosting framework	High performance, speed, and accuracy on structured data; regularization to prevent overfitting; parallel processing; widely used in Kaggle competitions.
	LightGBM	Fast, distributed, high-performance Gradient Boosting framework	Focus on efficiency and scalability, especially for large datasets; uses histogram-based algorithms for faster computation and less memory; developed by Microsoft.
	CatBoost	Gradient Boosting library with native categorical feature handling	Automatically handles categorical features without extensive preprocessing; reduces overfitting through ordered boosting; developed by Yandex.
Specialized NLP	Hugging Face Transformers	Provides pre-trained models for Natural Language Processing (NLP)	Access to thousands of state-of-the-art Transformer models (BERT, GPT, T5); simplifies advanced NLP tasks like text classification, sentiment analysis, and generation.
Computer Vision	OpenCV	Library of programming functions mainly aimed at real-time computer vision	Image and video processing; feature detection; object detection; often integrated with deep learning models for complex vision tasks.

Python Implementation and practical example: Predicting House Prices

1. Import Libraries

- `pandas`: For data handling.
- `sklearn.model_selection.train_test_split`: Split data into training/testing sets.
- `sklearn.linear_model.LinearRegression`: The prediction model.
- `sklearn.metrics.mean_squared_error`: Evaluate model accuracy.

2. Load Dataset

- `data = pd.read_csv("housing_data.csv")`: Load house data from a CSV file.

3. Feature Engineering

- `data['rooms_per_person'] = ...`: Create a new feature from existing ones (e.g., `rooms_per_person`).

4. Split Data

- `X = ..., y = ...`: Define features (input) and target (output).
- `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)`: Divide data into 80% for training and 20% for testing.

5. Train Model

- `model = LinearRegression()`: Initialize the linear regression model.
- `model.fit(X_train, y_train)` : Train the model using the training data.

6. Make Predictions

- `y_pred = model.predict(X_test)`: Use the trained model to predict values on the unseen test data.

7. Evaluate Model

- `mse = mean_squared_error(y_test, y_pred)`: Calculate how well the model performed using Mean Squared Error.
- `print(...)` : Display the evaluation result.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

# Load the dataset
data = pd.read_csv("housing_data.csv")

# Feature engineering (e.g., creating new features)
data['rooms_per_person'] = data['total_rooms'] / data['total_bedrooms']

# Split into training and testing sets
X = data[['total_rooms', 'total_bedrooms', 'rooms_per_person']]
y = data['median_house_value']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Create and train a linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
```

References

- [1] Tom Mitchell, Machine Learning. McGraw-Hill, 1997.
- [2] R. O. Duda, P. E. Hart, D. G. Stork, Pattern Classification, 2nd edition, John Wiley & Sons, 2000.
- [3] C. M. Bishop, “Pattern Recognition and Machine Learning”, Springer, 2006.