

## Practice exercise 4

### Multilayer Perceptron Regression Exercise: Predicting House Prices

**Problem Statement:** Given a dataset of houses with features like square footage, number of bedrooms, and location, predict the house price using a Multilayer Perceptron (MLP) model.

**Dataset:** You can use a popular dataset like the Boston Housing dataset, which is available in scikit-learn.

**Task:**

**1. Import Necessary Libraries.**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

**2. Load and Preprocess the Data:**

- Load the Boston Housing dataset.
- Split the data into features (X) and target variable (y).
- Scale the features using StandardScaler.
- Split the data into training and testing sets.

**3. Create a Multilayer Perceptron Model:**

- Define a Sequential model.
- Add Dense layers with appropriate activation functions (e.g., ReLU for hidden layers and linear for the output layer).
- Compile the model with an optimizer (e.g., Adam) and loss function (e.g., Mean Squared Error mse).

**4. Train the Model:**

- Fit the model to the training data.
- Monitor the training process using metrics like Mean Squared Error (MSE) and Mean Absolute Error (MAE).

**5. Evaluate the Model:**

- Evaluate the model on the testing set using the same metrics.
- Make predictions on the testing set and compare them to the actual value

**Code Implementation:**

```
from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split

# Load the Boston Housing dataset
boston = load_boston()
```

```

X = boston.data
y = boston.target

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Create the MLP model
model = Sequential([Dense(32, activation='relu', input_dim=X_train.shape[1]), Dense(16,
activation='relu'), Dense(1)])

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Train the model
model.fit(X_train_scaled, y_train, epochs=100, batch_size=32,
validation_data=(X_test_scaled, y_test))

# Evaluate the model
test_loss = model.evaluate(X_test_scaled, y_test)
print('Test Loss:', test_loss)

# Make predictions
y_pred = model.predict(X_test_scaled)

# Visualize the predictions
plt.scatter(y_test, y_pred)
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Actual vs. Predicted House Prices')
plt.show()

```

## 6. Compare the performance metrics and analyze the results.

- Predict the house price using various regression models: Linear Regression, Decision Tree Regression, Random Forest Regression on the testing set
- Calculate MSE and R-squared scores for each model.
- Compare the performance metrics with obtained predicted results of MLP model and analyze the results.